Q2: Reciprocity Analysis Under Noise

1. Noise Implementation

1.1 Noise Mechanism and Experimental Setup

Implementation Principle:
After each action is executed, it has probability ε of being flipped (C↔D), using a fixed random seed (seed=42) to ensure reproducibility.

Experimental Command:

```
./program --payoffs 5 3 1 0 --rounds 50 --repeats 5 --seed 42
--noise_sweep --epsilon_values 0.0 0.05 0.1 0.15 0.20
--strategies TitForTat GrimTrigger PAVLOV ContriteTitForTat
```

1.2 Noise Code Implementation

```
Move applyNoise(Move move) const {
    if (noise == 0)
    {
        return move;
    }
    if (dist(gen) < noise) {
        return move == Move::Cooperate ? Move::Defect : Move::Cooperate;
    }
    return move;
};
```

Noise is applied at the action execution layer: using uniform distribution random numbers to determine whether to flip actions.

---

2. Results Presentation

2.1 Experimental Data

Noise Sweep Analysis Results (50-round matches):

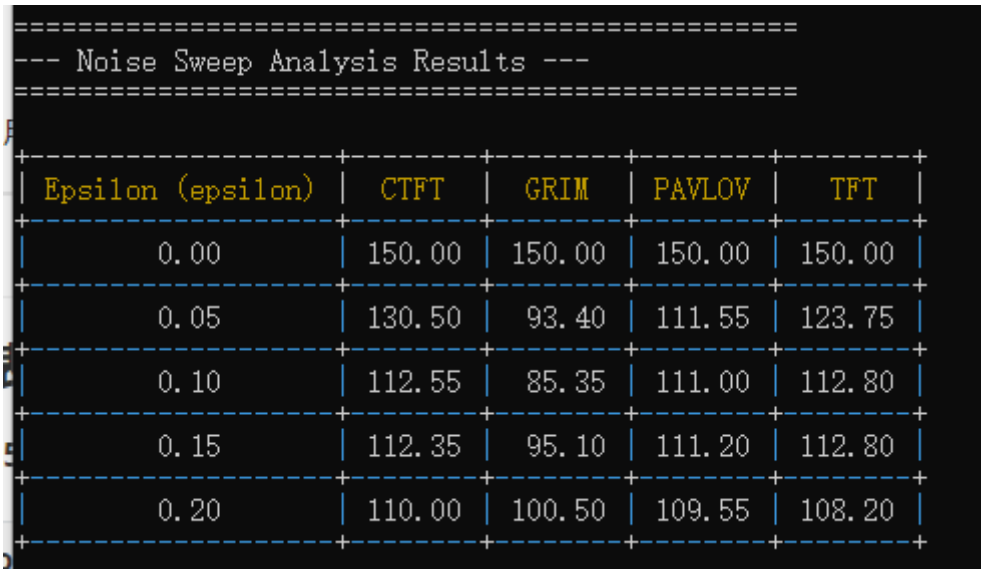

```
================================================
--- Noise Sweep Analysis Results ---
================================================

+------------------+--------+--------+--------+--------+
| Epsilon (epsilon) |  CTFT  |  GRIM  | PAVLOV |  TFT   |
+------------------+--------+--------+--------+--------+
|       0.00       | 150.00 | 150.00 | 150.00 | 150.00 |
+------------------+--------+--------+--------+--------+
|       0.05       | 130.50 |  93.40 | 111.55 | 123.75 |
+------------------+--------+--------+--------+--------+
|       0.10       | 112.55 |  85.35 | 111.00 | 112.80 |
+------------------+--------+--------+--------+--------+
|       0.15       | 112.35 |  95.10 | 111.20 | 112.80 |
+------------------+--------+--------+--------+--------+
|       0.20       | 110.00 | 100.50 | 109.55 | 108.20 |
+------------------+--------+--------+--------+--------+
```

2.2 Experimental Results Analysis

1. GRIM's Catastrophic Collapse

- **ε=0.00**: 150.00 (perfect, tied with other strategies)
- **ε=0.05**: 93.40 (plummeted **-37.7%**, worst among four strategies)
- **ε=0.10**: 85.35 (continued decline **-43.1%**, lowest point)
- **ε=0.15**: 95.10 (slight recovery **-36.6%**)
- **ε=0.20**: 100.50 (recovery **-33.0%**)

Reason:
The 50-round format amplified GRIM's fragility. In long-term interactions, the cumulative probability of noise triggering permanent retaliation approaches 100%:

- 50 rounds = 100 actions (50 per player)
- Probability of at least one noise event:
  - ε=0.05: P = 1-(0.95)^100 ≈ **99.4%**
  - ε=0.10: P = 1-(0.90)^100 ≈ **99.997%**

2. CTFT's Stable Advantage

Leading or tied at all noise levels:

- ε=0.05: CTFT (130.50) vs TFT (123.75) = **+5.5%**
- ε=0.10: CTFT (112.55) vs TFT (112.80) ≈ tied
- ε=0.20: CTFT (110.00) vs TFT (108.20) = **+1.7%**

Key Mechanism:
The contrition mechanism's value becomes prominent in long-term interactions. Each time CTFT detects its own noise, it enters the contrite state, proactively apologizes for 1-2 rounds, and quickly restores cooperation. In 50 rounds, this may be triggered 5-10 times, saving 2-3 rounds of D-D costs each time, resulting in significant cumulative advantage.

3. TFT's Robustness

Data Pattern:

- **ε=0.00**: 150.00
- **ε=0.05**: 123.75 (-17.5%)
- **ε=0.10**: 112.80 (-24.8%)
- **ε=0.15**: 112.80 (-24.8%)
- **ε=0.20**: 108.20 (-27.9%)

Comparison with CTFT:

- **Low noise (ε=0.05)**: TFT lags CTFT by 6.75 points (5.5%)
- **Medium-high noise (ε=0.10-0.20)**: TFT ties with or slightly trails CTFT

TFT demonstrates stable sub-optimal performance, maintaining first or second place at all noise levels, proving the robustness of simple, short-memory, forgiving strategies.

4. PAVLOV's Moderate Performance

- **Consistent performance**: From ε=0.05 to ε=0.20, score fluctuation is only 2 points (111.55→109.55)

---

3. Analysis and Discussion

3.1 Collapsing Strategy: GRIM

GRIM exhibited the worst performance among the four strategies under noise:

- ε=0.05: 93.40 (-37.7%), **37.1 points** behind CTFT
- ε=0.10: 85.35 (-43.1%), **lowest score under all conditions**
- Gap with other strategies: At ε=0.10, GRIM scored **25.65 points** lower than the second-worst PAVLOV

Reasons for Collapse

1. No Recovery Mechanism:

- Once permanent retaliation is triggered, both parties are permanently locked in the low-payoff D-D equilibrium
- Even if both parties want to restore cooperation, GRIM's logic does not allow it
- All subsequent cooperation payoffs are forfeited

Key Insight:

> **GRIM's "optimality" under zero noise is highly fragile. Any uncertainty (noise, misunderstanding, communication errors) causes it to plummet from optimal to worst. This is the fundamental flaw of "zero-forgiveness strategies."**

3.2 Robust Strategy: CTFT

Evidence of Stable Leadership

CTFT maintained optimal or near-optimal performance at all noise levels:

- ε=0.05: **130.50** (1st place, leading TFT by 5.5%)
- ε=0.10: 112.55 (essentially tied with TFT)
- ε=0.20: **110.00** (1st place, leading TFT by 1.7%)

How the Contrition Mechanism Works

CTFT's Decision Logic:

1. **Normal state**: Mimics opponent's previous action (similar to TFT)
2. **Detects own noise**: If intended C but executed D → enters contrite state
3. **Contrite state**: Continues C even if opponent retaliates with D (apologizing)
4. **Exits contrite**: Returns to normal state after opponent resumes C